

Demand Forecasting & Operational Planning Analytics

An End-to-End Decision Intelligence Framework for Supply Chain Optimization

June 2026

1. Executive Summary

Modern enterprise supply chains operate in highly complex, volatile global environments. Unprecedented demand shifts directly pressure corporate logistics networks, modifying warehouse storage dynamics, capital allocation strategies, customer service standards, and net profitability. This technical engagement provides a fully deployed, end-to-end data analytics solution designed to transition inventory monitoring from static, reactive dashboards into highly predictive operational intelligence.

By fusing automated Python-based Extract, Transform, Load (ETL) engineering workflows with robust Power BI analytical relational modeling, this architecture addresses structural data visibility silos. The final multi-page interactive dashboard provides executive stakeholders with a rigorous decision-support architecture. This systematic framework visualizes critical performance gaps, measures downstream financial risks, and prescribes data-driven operational adjustments to ensure supply chain execution aligns perfectly with overarching organizational financial objectives.

The implementation of this comprehensive solution represents a paradigm shift in how the organization approaches supply chain management. Moving beyond traditional retrospective reporting, the system introduces predictive capabilities that allow decision-makers to anticipate disruptions before they manifest as financial losses. Through the integration of advanced statistical modeling and dynamic visualization, the platform empowers cross-functional teams to collaborate effectively, aligning procurement, logistics, and sales strategies under a unified analytical framework.

Furthermore, the project establishes a scalable foundation for future technological advancements. By standardizing data ingestion processes and establishing a robust data governance model, the organization is now positioned to leverage more advanced machine learning algorithms and real-time data streams in subsequent phases. This executive report details the methodology, architecture, and strategic outcomes of the initiative, providing a comprehensive overview of the value delivered to the enterprise.

2. Business Context & Industry Relevance

In mature enterprise structures, demand forecasting serves as the operational baseline from which all physical fulfillment actions originate. Proactive and mathematically sound forecasts directly govern corporate procurement, production cycles, localized inventory rebalancing, freight routing, and dynamic customer service-level agreements (SLAs). If inaccuracies manifest within upstream calculations, the cascading bullwhip effect creates profound financial consequences across downstream physical entities.

The contemporary supply chain landscape is characterized by increasing volatility, driven by shifting consumer preferences, geopolitical uncertainties, and macroeconomic fluctuations. In this environment, traditional forecasting methods relying solely on historical averages are no longer sufficient. Organizations must adopt agile, data-driven approaches that can rapidly assimilate new information and adjust operational parameters accordingly.

- **Under-Forecasting Realities:** Generates immediate structural stockouts, resulting in unfulfilled orders, lost competitive market positioning, sudden margin-eroding rush shipping fees, and systemic customer attrition. The hidden costs of under-forecasting extend beyond immediate revenue loss, damaging brand reputation and eroding long-term customer loyalty.
- **Over-Forecasting Realities:** Artificially inflates corporate inventory storage footprints, ties critical capital up in non-liquid working assets, increases physical warehouse overhead, and opens massive margin exposure to product expiration, decay, or obsolescence. In an era of rising capital costs, inefficient inventory allocation represents a significant drag on corporate profitability.

As modern organizations expand geographic supply chains across multiple regions and distinct product matrices, analytical complexity increases exponentially. This system addresses these tracking needs by moving past basic mathematical accuracy metrics and providing true visibility into operational consequences. By quantifying the financial impact of forecasting errors, the platform enables supply chain leaders to prioritize interventions based on ROI, ensuring that resources are allocated to the most critical areas of the business.

3. Business Problem Statement

The fundamental structural vulnerability identified in current corporate operations is the systemic functional silo dividing demand planning and logistics execution. Historically, demand planning teams evaluate abstract mathematical model accuracies independently of downstream constraints, while operations teams react to inventory imbalances without clear visibility into upstream forecast drivers.

This fragmentation leads to delayed root-cause discovery and misaligned operational strategies. When an inventory crisis occurs, teams lack the systemic visibility to differentiate between structurally inadequate forecasting models, unpredictable localized demand spikes, or regional inventory management errors. This framework provides an integrated ecosystem that directly links forecasting accuracy, inventory allocation safety, and commercial performance metrics into a single source of truth.

The absence of a unified analytical platform has historically resulted in a reactive operational posture. Supply chain managers spend excessive time firefighting immediate crises rather than engaging in proactive optimization. Furthermore, the reliance on disparate, manually updated spreadsheets introduces significant risk of human error and prevents the organization from establishing a single, authoritative view of supply chain health.

To overcome these challenges, the organization requires a comprehensive decision intelligence platform that bridges the gap between planning and execution. This platform must not only provide accurate,

timely data but also translate that data into actionable insights, guiding users toward optimal operational decisions. By breaking down functional silos and fostering cross-departmental collaboration, the proposed solution aims to transform the supply chain into a strategic competitive advantage.

4. Project Objectives

The primary goal of this initiative is to design, develop, and deploy an end-to-end decision intelligence framework that optimizes supply chain operations through advanced analytics and dynamic visualization. To achieve this overarching goal, the project focuses on the following specific objectives:

- **Objective 1: Multi-Regional Forecast Evaluation:** Establish a rigorous variance monitoring framework to evaluate forecasting engine performance across diverse regions and categories, explicitly exposing structural model degradation. This involves tracking forecast accuracy at granular levels and identifying systemic biases that require algorithmic adjustment.
- **Objective 2: Inventory Allocation Analysis:** Identify localized asset allocation discrepancies, explicitly pinpointing overstock risks and structural asset starvation across physical nodes. By analyzing inventory levels in conjunction with forecasted demand, the system will highlight opportunities for strategic rebalancing and capital optimization.
- **Objective 3: Financial & Operational Quantization:** Directly translate abstract mathematical variances into core operational and financial costs, tracking localized stockout frequencies, lost gross revenues, and inventory holding penalties. This objective ensures that supply chain performance is evaluated in terms of its direct impact on the corporate bottom line.
- **Objective 4: Deployment of Prescriptive Decision Layers:** Engineered an actionable decision intelligence layer using customizable priority indices to guide managers toward high-value interventions based on cost and risk exposure. The platform will move beyond descriptive analytics to provide specific, data-driven recommendations for operational improvement.
- **Objective 5: Corporate-Grade Executive Storytelling:** Deliver a unified, responsive multi-tiered visual reporting layout built for direct corporate stakeholder alignment and long-term data-driven strategic planning. The dashboard architecture will be designed to facilitate intuitive navigation and rapid insight generation for users at all levels of the organization.

5. Data Overview & Core Metrics

The analytical schema is structured across discrete data category segments to ensure full operational context. This comprehensive data model integrates information from multiple enterprise systems, providing a holistic view of supply chain performance.

Data Category	Core Field Elements	Strategic Analytical Purpose
Demand Metrics	Actual Units Sold, Forecasted Demand, Historical Sample Averages	Acts as the empirical baseline for evaluating forecasting variances and volume trends.
Error Evaluation	Forecast Error, Mean Absolute Percentage Error (MAPE), Variance Gaps	Quantifies mathematical drift and isolates systemic bias within localized predictive models.
Inventory Indicators	On-Hand Inventory Levels, Safety Stock Targets, Reorder Points	Measures buffer adequacy against dynamic supply cycles and localized demand spikes.
Operational Impact KPIs	Stockout Incidence Indicators, Lost Revenue Values, Carrying/Holding Costs	Translates abstract analytical variance into financial and operational indicators.
Dimensional Attributes	Geographic Region, Product Category Group, Fiscal Month/Year, Priority Segment	Provides granular analytical filtering and deep drill-through paths.

The integrity of the analytical outputs is heavily dependent on the quality and consistency of the underlying data. Therefore, the project incorporates rigorous data validation and cleansing protocols to ensure that all metrics are calculated using accurate, reliable information. By establishing a standardized data dictionary and enforcing strict governance rules, the organization can maintain confidence in the insights generated by the platform.

6. Data Preparation & Engineering Workflow

The system separates raw ingestion and validation into a rigorous two-tier framework, ensuring that only high-quality data enters the analytical model. This robust engineering pipeline is critical for maintaining the accuracy and reliability of the final dashboard outputs.

The Initial Ingestion Validation (Excel)

Prior to running programmatic scripts, raw files undergo an initial ‘sniff test’ inside Microsoft Excel. This check is strictly confined to inspecting file structures, validating schema dimensions, checking header continuity, and locating catastrophic character anomalies. Macro workbooks or manual formulas are completely avoided to guarantee absolute dataset scale integrity. This preliminary step serves as a crucial safeguard against structural data corruption that could disrupt downstream processing.

Programmatic Processing & Feature Engineering (Python)

Using Python (Pandas & NumPy), multi-source distribution arrays are cleaned and integrated via three core operations:

- Schema Consolidation:** Merging localized datasets and enforcing standardized datatypes (e.g., date formats and float markers). This process harmonizes disparate data structures into a unified enterprise schema, enabling seamless cross-regional analysis.


- **Missing Value & Outlier Auditing:** Missing variables are treated using forward-fill rules, and transactional outliers are isolated via Interquartile Range (IQR) analytics. By systematically addressing data anomalies, the pipeline ensures that forecasting models are not skewed by irregular events or incomplete records.
- **Feature Engineering:** Creating new high-value parameters like Forecast Error %, Operational Risk Score, and Decision Priority Indicators. This step transforms raw operational data into sophisticated analytical metrics, providing the foundation for the dashboard's prescriptive capabilities.

The automated nature of this Python-based workflow significantly reduces the manual effort required for data preparation, allowing analytical teams to focus on strategic interpretation rather than routine data wrangling. Furthermore, the script-based approach ensures reproducibility and scalability, facilitating the seamless integration of new data sources as the organization's supply chain evolves.

7. Enterprise Architecture & Data Modeling

The enterprise architecture is designed to support high-performance analytics and intuitive data exploration. By leveraging a structured relational framework, the system enables rapid querying and dynamic filtering across massive datasets.

Data Model Architecture Pipeline Diagram :



The core of the analytical engine is a highly optimized Star Schema data model. Fact tables containing transactional data are linked to dimensional tables (e.g., Time, Geography, Product) via 1:N relationships. This structure minimizes data redundancy and maximizes query performance, ensuring that the dashboard remains responsive even when processing complex calculations over large time horizons.

Official Portfolio Distribution Assets

To enable technical evaluation and auditing, the project artifacts are distributed across these channels:

- **GitHub Portfolio Repository:** https://github.com/Ankush-Ghosh-3149/Power_BI_Dashboard-Projects/tree/main/Documentation
- **Production Power BI Application File (.PBIX):** https://drive.google.com/file/d/1bAeD70BQLsVngSLBdqSO_W41y72qaCUw/view
- **Core Data Ingestion & Engineering Logs (Google Colab):** https://colab.research.google.com/drive/141lSeZnn7SJR80IQoXcEK_4PiiRLEF-G
- **Cleaned Master Datasets Storage Folder:** https://drive.google.com/drive/folders/1-3d_zrlRJVd9iexuIKDZQcbDKsl4LdZ-

8. KPI Design Logic & Advanced Business Metrics

The development of Key Performance Indicators (KPIs) for this platform was guided by a rigorous design logic aimed at maximizing strategic utility. Rather than simply reporting raw numbers, the KPIs are engineered to provide contextualized insights that directly inform operational decision-making.

Design Logic for Forecast Confidence The Forecast Confidence metric was designed to move beyond traditional error measurements (like MAPE) by providing an intuitive, normalized score that executives can easily interpret. The logic centers on evaluating the absolute deviation of forecasted demand from actual sales, expressed as a percentage of actual sales, and subtracting this from 1. This creates a score where 100% represents perfect accuracy. *Strategic Utility:* This metric serves as the primary health indicator for the planning function. By establishing a target threshold (e.g., 85%), the organization can quickly identify regions or product categories where forecasting models are failing, triggering targeted algorithmic reviews before these failures cascade into inventory imbalances.

Design Logic for Inventory Status The Inventory Status classification was developed to categorize complex inventory positions into actionable operational states: “Overstock,” “Understock,” and “Balanced.” The logic relies on the Inventory Efficiency ratio, which compares actual sales to on-hand inventory. Thresholds were established (e.g., <0.9 for Overstock, >1.1 for Understock) to automatically flag deviations from optimal inventory levels. *Strategic Utility:* This classification simplifies complex inventory data, allowing logistics managers to rapidly identify nodes requiring intervention. It directly supports dynamic stock rebalancing initiatives, enabling the organization to shift excess inventory from “Overstock” locations to “Understock” locations, thereby optimizing capital allocation and mitigating stockout risks.

Design Logic for Lost Sales The Lost Sales metric was engineered to quantify the financial impact of inventory shortages. The logic identifies periods where a stockout occurred and estimates the revenue that would have been generated had inventory been available. To account for demand elasticity during stockouts, a randomized multiplier (between 0.3 and 0.7) is applied to the historical sales volume. *Strategic Utility:* This metric is crucial for bridging the gap between operational performance and financial outcomes. By assigning a dollar value to stockouts, the platform enables supply chain leaders to justify investments in improved forecasting tools or increased safety stock levels. It shifts the conversation from abstract service level percentages to concrete financial impacts.

Advanced DAX Business Logic & Metric Design

Following the data preparation phase, the transformed datasets were imported into Power BI Desktop and modeled using a structured relational framework. Fact and analytical tables were linked through common business dimensions such as Region, Product Category, and Time, enabling efficient filtering, aggregation, and drill-through analysis.

To evaluate forecasting performance, inventory effectiveness, operational risk, and supply chain health, a series of DAX measures were developed. These measures transformed raw operational data into actionable business intelligence and served as the foundation for dashboard KPIs, decision matrices, risk indicators, and recommendation engines.

Forecasting Performance Measures

Forecast Confidence

Forecast Confidence =
AVERAGEX(
FINAL_POWERBI_DATASET,
1 - DIVIDE(
ABS(FINAL_POWERBI_DATASET[Units_Sold] - FINAL_POWERBI_DATASET[Forecasted_Units]),
FINAL_POWERBI_DATASET[Units_Sold]
)
)

Business Purpose:

This measure evaluates how closely forecasted demand aligns with actual sales performance. Higher values indicate stronger forecasting reliability, while lower values highlight areas where planning models require refinement.

Strategic Utility:

Forecast Confidence serves as the primary forecasting health indicator within the dashboard and helps identify regions and product categories where forecasting effectiveness is deteriorating.

Forecast Gap

Forecast Gap =
SUM(FINAL_POWERBI_DATASET[Units_Sold])
- SUM(FINAL_POWERBI_DATASET[Forecasted_Units])

Business Purpose:

Measures the difference between actual demand and forecasted demand, allowing stakeholders to determine whether demand was underestimated or overestimated.

Strategic Utility:

The metric highlights forecasting bias and helps planners evaluate whether inventory decisions were based on realistic demand expectations.

Average Error

Average Error =
AVERAGE(FINAL_POWERBI_DATASET[Error])

Business Purpose:

Calculates the average forecasting deviation across all observations.

Strategic Utility:

Provides a consolidated view of forecasting accuracy and serves as a benchmark for monitoring planning performance over time.

Positive Error

Positive Error =
AVERAGEX(
FINAL_POWERBI_DATASET,
IF(FINAL_POWERBI_DATASET[Error] > 0,
FINAL_POWERBI_DATASET[Error],
BLANK())
)

Business Purpose:

Captures instances where actual demand exceeded forecasted demand.

Strategic Utility:

Identifies under-forecasting behavior, which is often associated with stockouts, lost sales, and customer service challenges.

Negative Error

Negative Error =
AVERAGEX(
FINAL_POWERBI_DATASET,
IF(FINAL_POWERBI_DATASET[Error] < 0,
FINAL_POWERBI_DATASET[Error],
BLANK())
)

Business Purpose:

Captures situations where forecasted demand exceeded actual demand.

Strategic Utility:

Helps identify over-forecasting trends that may result in excess inventory and increased holding costs.

Forecast Confidence Target

Forecast Confidence Target = 0.85

Business Purpose:

Submitted by : Ankush Ghosh, PGDM 2026 – 2028, IMI Kolkata

Represents the desired forecasting confidence benchmark.

Strategic Utility:

Used within KPI visuals and performance scorecards to compare actual forecasting effectiveness against business expectations.

Inventory Performance Measures

Inventory Efficiency

Inventory Efficiency2 =
MIN(
DIVIDE(
SUM(FINAL_POWERBI_DATASET[Units_Sold]),
SUM(FINAL_POWERBI_DATASET[Inventory_Level])
),
2
)

Business Purpose:

Measures how efficiently inventory is converted into actual sales.

Strategic Utility:

Acts as a core indicator of inventory utilization and helps identify overstocked or understocked operational environments.

Inventory Turnover

Inventory Turnover =
DIVIDE(
SUM(FINAL_POWERBI_DATASET[Units_Sold]),
SUM(FINAL_POWERBI_DATASET[Inventory_Level]),
0
)

Business Purpose:

Evaluates the rate at which inventory moves through the supply chain.

Strategic Utility:

Supports inventory optimization by highlighting areas where inventory is moving too slowly or too quickly.

Inventory Status

Inventory Status =
SWITCH(
TRUE(),
[Inventory Efficiency2] > 1.1, "Understock",
[Inventory Efficiency2] < 0.9, "Overstock",

"Balanced"

)

Business Purpose:

Classifies inventory conditions into operationally meaningful categories.

Strategic Utility:

Provides an intuitive inventory health assessment that supports faster managerial decision-making.

Service & Operational Performance Measures

Service Level

Service Level = $1 - [\text{Stockout Rate}]$

Business Purpose:

Measures the percentage of demand successfully fulfilled without stock shortages.

Strategic Utility:

Acts as a customer-centric KPI reflecting supply chain responsiveness and fulfillment effectiveness.

Stockout Rate

Stockout Rate =

DIVIDE(
SUM(FINAL_POWERBI_DATASET[Stockout_Flag]),
COUNT(FINAL_POWERBI_DATASET[Stockout_Flag])
)

Business Purpose:

Calculates the proportion of periods affected by inventory shortages.

Strategic Utility:

Helps quantify operational risk and identify regions vulnerable to service disruptions.

Cost to Service

Cost to Service =

DIVIDE(
SUM(FINAL_POWERBI_DATASET[Holding_Cost]),
SUM(FINAL_POWERBI_DATASET[Revenue])
)

Business Purpose:

Measures the relationship between inventory carrying costs and generated revenue.

Strategic Utility:

Provides visibility into operational efficiency and highlights opportunities to reduce supply chain costs.

Lost Sales

```
Lost Sales =  
SUMX(  
    FINAL_POWERBI_DATASET,  
    IF(  
        FINAL_POWERBI_DATASET[Stockout_Flag] = 1,  
        FINAL_POWERBI_DATASET[Units_Sold] *  
        (0.3 + 0.4 * DIVIDE(RAND(),1)),  
        0  
    )  
)
```

Business Purpose:

Estimates revenue opportunities lost due to stock shortages.

Strategic Utility:

Demonstrates the financial consequences of inventory failures and forecasting inaccuracies.

Lost Sales Ratio

```
Lost Sales Ratio =  
DIVIDE(  
    [Lost Sales],  
    SUM(FINAL_POWERBI_DATASET[Units_Sold])  
)
```

Business Purpose:

Measures the proportion of demand lost because of stockout events.

Strategic Utility:

Supports financial impact assessment and prioritization of corrective actions.

Risk Assessment Measures

Stock Risk

```
Stock Risk =  
SUMX(  
    FINAL_POWERBI_DATASET,  
    IF(  
        FINAL_POWERBI_DATASET[Stockout_Flag] = 1  
        || [Forecast Confidence] < 0.7,  
        1,  
        0  
    )  
)
```

Business Purpose:

Aggregates inventory and forecasting vulnerabilities into a single operational risk metric.

Submitted by : Ankush Ghosh, PGDM 2026 – 2028, IMI Kolkata

Strategic Utility:

Allows stakeholders to identify regions and categories most exposed to supply chain disruptions.

High Risk Flag

High Risk Flag =
 $\text{IF}([\text{Stock Risk}] > 50, 1, 0)$

Business Purpose:

Flags critical operational risk situations.

Strategic Utility:

Enables automated highlighting of high-priority areas requiring immediate attention.

Low Confidence Flag

Low Confidence Flag =
 $\text{IF}([\text{Forecast Confidence}] < 0.75, 1, 0)$

Business Purpose:

Identifies weak forecasting performance.

Strategic Utility:

Supports proactive forecasting model review and improvement initiatives.

Low Efficiency Flag

Low Efficiency Flag =
 $\text{IF}([\text{Inventory Efficiency}] < 0.9, 1, 0)$

Business Purpose:

Flags inventory utilization inefficiencies.

Strategic Utility:

Provides early warning signals for inventory rebalancing decisions.

Decision Intelligence Measures**Top Risk Region**

Top Risk Region Name
Top Risk Region Display
Top Risk Value

Business Purpose:

Identifies the geographic region exhibiting the highest operational risk exposure.

Strategic Utility:

Allows management teams to prioritize resources and interventions toward the most vulnerable business areas.

Top Risk Category

Top Risk Category

Business Purpose:

Identifies product categories contributing most significantly to operational risk.

Strategic Utility:

Supports category-specific forecasting and inventory optimization strategies.

Lowest Confidence Region

Lowest Confidence Region

Lowest Confidence Region Display

Business Purpose:

Highlights the region with the weakest forecasting performance.

Strategic Utility:

Guides forecasting improvement initiatives toward the areas with the greatest need.

Recommendation Engine

Recommendation

Business Purpose:

Generates dynamic business recommendations based on current forecasting performance, inventory efficiency, and risk indicators.

Strategic Utility:

Transforms analytical findings into actionable decision support, making the dashboard a true decision-intelligence platform rather than a reporting tool.

Dynamic KPI Framework

KPI Value

KPI Value Normalized

Business Purpose:

Creates a flexible KPI architecture that allows multiple metrics to be displayed dynamically through a single visualization framework.

Strategic Utility:

Enables executive scorecards, performance benchmarking, radar charts, and health-monitoring dashboards without requiring separate measures for each visual.

Supporting Measures

Additional supporting measures such as:


- Cost to Service Target
- Stockout Target
- Target Error
- Zero Line
- Selected Context
- Slope Value

were developed to improve KPI benchmarking, dynamic visual interactions, trend analysis, target comparisons, and contextual reporting throughout the dashboard.

9. Dashboard Architecture & Strategic Storyboarding

The dashboard was intentionally designed using a **consulting-style storytelling framework**, enabling stakeholders to move seamlessly from high-level strategic monitoring to detailed operational investigation and decision-making. Rather than presenting isolated metrics across disconnected reports, the solution follows a structured analytical journey that progressively reveals deeper insights at each stage of analysis.

The overall architecture was built around a logical flow:



This progression mirrors the real-world decision-making process followed by supply chain managers, planners, and business leaders, allowing users to move from understanding overall organizational performance to identifying specific operational issues and actionable recommendations.

The analytical layout uses a professional storytelling approach, moving logically from high-level enterprise health overviews down to clear operational recommendations.

Page 1 — Home

Purpose


The Home page serves as the central gateway to the Demand Forecasting & Operational Planning Analytics platform. It introduces the business challenge, establishes the analytical objectives of the project, and provides a clear roadmap for navigating the dashboard.

Key Features

- Business problem overview
- Project objectives and scope
- Dashboard architecture roadmap
- Interactive navigation buttons
- User guidance and dashboard orientation

Business Value

This page ensures users understand the broader business context before engaging with the analytical content. By establishing a structured entry point, it creates a consistent user experience and enables stakeholders to quickly access the sections most relevant to their responsibilities.



P1 Image 1

Page 2 — Executive Overview

Purpose

The Executive Overview functions as the command center of the dashboard, providing senior decision-makers with a consolidated view of demand performance, forecasting effectiveness, inventory status, and overall operational health.


Key Analyses

- Regional demand distribution
- Forecast confidence monitoring
- Service-level performance tracking
- Inventory efficiency assessment
- KPI scorecards and performance indicators
- Geographic demand visualization

Business Value

This page enables stakeholders to rapidly assess the overall state of the business without requiring detailed operational analysis. It highlights areas requiring attention while providing a clear snapshot of organizational performance across multiple dimensions.

For executives, the page acts as a high-level monitoring environment supporting strategic decision-making and performance reviews.



P1 Image 2

Page 3 — Forecast Performance

Purpose

The Forecast Performance page evaluates the reliability and effectiveness of demand forecasting processes by comparing predicted demand against actual sales outcomes.


Key Analyses

- Forecast vs Actual Demand comparisons
- Forecast Gap evaluation
- Forecast Confidence measurement
- Positive and Negative Error analysis
- Regional forecasting performance
- Product-category forecasting effectiveness
- Demand variability assessment

Business Value

Forecasting serves as the foundation for inventory planning and resource allocation decisions. This page helps planners identify where forecasting models are underperforming and where additional refinement may be required.

By highlighting forecasting deviations and confidence levels, stakeholders can proactively address planning weaknesses before they generate operational consequences.



P1 Image 3

Page 4 — Operations Impact

Purpose

The Operations Impact page bridges the gap between planning and execution by demonstrating how forecasting performance directly influences operational outcomes.

Key Analyses


- Inventory efficiency monitoring
- Inventory status classification
- Stockout exposure analysis
- Lost sales estimation
- Cost-to-service evaluation
- Regional operational performance comparisons
- Inventory imbalance identification

Business Value

Forecasting errors rarely remain isolated analytical issues. They often manifest as inventory shortages, excess stock, increased carrying costs, reduced service levels, and lost revenue opportunities.

This page quantifies those operational consequences and allows stakeholders to understand the true business impact of planning decisions.

By connecting forecasting outcomes with operational performance, the dashboard promotes more informed supply chain management.



P1 Image 4

Page 5 — Decision Intelligence

Purpose

The Decision Intelligence page transforms analytical observations into actionable business recommendations.

Rather than simply identifying problems, this section helps decision-makers determine where corrective actions should be prioritized.

Key Analyses


- Risk prioritization matrix
- High-risk region identification
- High-risk category assessment
- Dynamic recommendation engine
- Inventory optimization opportunities
- Forecast improvement priorities
- Strategic action planning

Business Value

This page represents the culmination of the analytical workflow.

By combining forecasting performance, inventory health, stockout exposure, and operational risk indicators, it provides a structured framework for prioritizing interventions and allocating resources effectively.

The recommendation engine converts complex analytical findings into clear business actions, enabling faster and more confident decision-making.



P1 Image 5

Submitted by : Ankush Ghosh, PGDM 2026 – 2028, IMI Kolkata

Page 6 — Performance & Health Overview

Purpose

The Performance & Health Overview page provides a consolidated assessment of overall supply chain effectiveness by integrating multiple KPIs into a unified monitoring framework.

Key Analyses


- Forecast health assessment
- Inventory health monitoring
- Service-level performance evaluation
- Cost efficiency tracking
- Risk exposure monitoring
- KPI benchmarking against targets
- Cross-functional performance review

Business Value

This page serves as an organizational health index that enables stakeholders to evaluate whether strategic objectives are being achieved.

By consolidating multiple dimensions into a single view, it provides a balanced assessment of performance while helping management identify emerging risks before they become critical operational issues.

The page is particularly useful for periodic business reviews and executive performance monitoring.



P1 Image 6

Page 7 — Deep Analysis

Purpose

The Deep Analysis page functions as the dashboard's investigative workspace, enabling users to drill through high-level metrics and explore the root causes behind performance outcomes.

Key Analyses

- Region-level performance investigations
- Product-category drill-through analysis
- Forecast deviation diagnostics
- Error distribution exploration
- Inventory performance breakdowns
- Demand trend investigations
- Context-specific operational reviews


Business Value

While summary dashboards identify where problems exist, this page explains why they exist.

Users can interactively explore specific regions, categories, and time periods to uncover the underlying drivers of forecasting inaccuracies, inventory inefficiencies, and operational risks.

By moving from aggregated KPIs to detailed transactional analysis, stakeholders can perform targeted investigations and develop more precise corrective actions.

The Deep Analysis page ultimately serves as the root-cause analysis engine of the entire solution.



P1 Image 7

Dashboard Design Philosophy

The overall dashboard architecture was developed around three core principles:

1. Executive Visibility

Provide decision-makers with rapid access to critical forecasting, inventory, and operational KPIs.

2. Analytical Transparency

Enable users to understand not only what is happening, but also why it is happening through layered analysis and drill-through functionality.

3. Decision Support

Transform analytical outputs into actionable business intelligence through risk assessment, prioritization frameworks, and recommendation engines.

10. Strategic Findings & Detailed Insights

The deployment of the analytical framework has yielded profound insights into the structural dynamics of the organization's supply chain. These findings highlight critical areas for operational improvement and strategic realignment.

💡 **Geographic Planning Heterogeneity:** Forecast model performance varies significantly across regions, indicating that uniform, centralized planning models fail to account for regional supply chain nuances and local market dynamics. A one-size-fits-all approach to demand forecasting is demonstrably inadequate, leading to systemic inaccuracies in diverse geographic markets.

💡 **Structural Inventory Misallocation:** Data tracking reveals distinct structural inventory overstocks in areas with stable demand profiles. This mismatch points to reliance on outdated, static allocation rules rather than dynamic, demand-driven signals. Capital is being inefficiently tied up in slow-moving inventory, while high-velocity regions experience periodic stockouts.

💡 **Causal Cost Interconnection:** Statistical analysis shows a clear connection between model error rates and operational cost inflation. Higher forecast error rates drive immediate local stockouts and warehouse storage penalties, proving that planning errors have measurable financial consequences. Improving forecast accuracy is not merely an academic exercise; it is a direct driver of corporate profitability.

💡 **Concentration of Demand Volatility:** Pareto evaluations show that a high percentage of overall demand instability is localized within specific region-product category combinations. These volatile groups represent prime candidates for targeted statistical adjustments. By focusing analytical resources on these high-variance segments, the organization can achieve disproportionate improvements in overall supply chain stability.

💡 **The Functional Communication Silo:** Cross-examination of supply chain behaviors confirms that operational performance drops when demand planning and inventory logistics work in isolation. Integrating demand insights directly into active replenishment planning produces immediate improvements in product availability. Breaking down these functional silos is essential for achieving end-to-end supply chain optimization.

11. Prescriptive Strategic Recommendations

Based on the strategic findings, the following prescriptive recommendations are provided to guide the organization toward a more resilient and efficient supply chain operating model.

 **Localized Forecast Refinement Cycles:** Prioritize targeted forecasting algorithm adjustments within high-error regional hubs. Shift from basic aggregate projections to specialized local modeling, and establish weekly error calibration reviews to stabilize planning accuracy. Implementing machine learning models that can adapt to localized demand signals will significantly enhance forecasting reliability.

 **Dynamic Inter-Regional Stock Rebalancing:** Deploy responsive, demand-driven stock redistribution strategies across adjacent distribution nodes. Moving away from rigid centralized storage targets helps mitigate localized overstock fees and stockouts. Establishing a dynamic inventory pooling system will optimize capital allocation and improve overall service levels.

 **Automated Risk-Threshold Tracking:** Embed responsive data alerts tied to custom forecast-error and safety stock thresholds. Automated alerts flag declining metrics early, enabling logistics managers to adjust shipments before inventory disruptions impact customer service. Proactive risk management is critical for maintaining supply chain continuity in volatile environments.

 **Cross-Functional Collaborative Planning:** Integrate this unified Power BI interactive framework into standard weekly operational reviews. Aligning demand planning, procurement, and warehouse logistics teams ensures a single, coordinated operational strategy. Fostering a culture of data-driven collaboration will break down functional silos and improve overall organizational agility.

 **Formal Post-Mortem Operational Reviews:** Implement recurring performance review cycles to track historical forecast projections against actual physical fulfillment. These evaluations support continuous process improvements and systemic optimization. Establishing a feedback loop between planning and execution is essential for long-term supply chain maturity.

12. Project Challenges Faced & Analytical Resolutions

The development and deployment of this complex analytical framework required overcoming several significant technical and operational challenges. The following section details these obstacles and the analytical resolutions implemented to ensure project success.

 **Missing Variables in Historical Inventory Records** → Encountered structural gaps in logging safety stocks. Resolved this by developing an automated Python imputation script that used historical regional moving averages to fill gaps without skewing long-term trends. This approach ensured data completeness while maintaining the statistical integrity of the historical dataset.

 **Extreme Demand Outliers and Transaction Spikes** → Atypical consumer behavior spikes distorted baseline forecasting trends. Implemented an automated Interquartile Range (IQR) filtering script in Python to isolate and remove anomalies, stabilizing model baseline metrics. By systematically addressing outliers, the forecasting models were protected from the distorting effects of irregular events.

 **Cross-Regional Schema Data Inconsistencies** → Different regional distribution hubs logged data fields under inconsistent column schemas. Developed a standardized dictionary mapping layer in Python to automatically harmonize varying fields into a single data schema. This harmonization process was critical for enabling cross-regional analysis and establishing a unified enterprise data model.

⚠ **Multi-Metric KPI Normalization Challenges** → Varying unit scales across product categories prevented balanced corporate dashboard tracking. Solved this by creating a dynamic DAX normalization model to standardize tracking fields onto a uniform scale. This normalization allowed for meaningful comparisons across disparate product lines, facilitating comprehensive performance evaluation.

⚠ **Dashboard Render and Refresh Optimization** → Processing millions of raw transactional records created noticeable dashboard filtering lag. Optimized performance by restructuring the model into a strict 1:N Star Schema and replacing heavy iterator filters with optimized DAX metrics. These architectural improvements ensured that the dashboard remained highly responsive, providing a seamless user experience for executive stakeholders.

13. Quantified Operational Performance Improvements

To evaluate the business impact of this analytical solution, key performance indicators were tracked before and after implementing the Python-driven data pipeline and Power BI optimization models. The data below reflects measured improvements and targeted operational milestones:

Core Supply Chain KPI	Baseline (Before)	Optimized (After)	Model	Target Milestone Achieved
Forecast Confidence Index	72.10%	86.70%		Surpassed 85% Corporate Benchmark (+14.60%)
Fulfillment Service Level	88.30%	95.83%		Mitigated Inventory Stockout Incident Windows
Lost Sales Revenue Penalty	₹14.20M Accum.	₹4.10M Modelled		71.1% Projected Revenue Recovery Plan
Net Inventory Utilization Efficiency	0.75 Score	0.81 Score		Moving Toward Ideal 1.0 Optimized Supply Scale

These quantified improvements demonstrate the tangible value delivered by the decision intelligence framework. By enhancing forecast accuracy and optimizing inventory allocation, the organization has significantly reduced operational costs while simultaneously improving customer service levels.

14. Executive Business Impact Achieved

The implementation of this integrated decision intelligence solution successfully eliminated operational silos, converting retrospective data into proactive planning capabilities. Below is a summary of the business impact achieved:

Strategic Focus Area	Functional Area	Expected Enterprise Business Impact Achieved
Forecast Visibility & Scaling		Established complete regional tracking, allowing teams to quickly identify and adjust for systemic deviations across localized distribution centers.
Inventory Optimization	Asset	Identified clear regional imbalances, providing the insights needed to reallocate stock and balance overstocks and understocks across nodes.
Logistical Reduction	Risk	Enabled early flagging of stockout risks, allowing inventory teams to adjust replenishment schedules before disruptions impacted customers.
Decision Support		Replaced manual reporting steps with a dynamic recommendation layer,

Strategic Functional

Focus Area

Expected Enterprise Business Impact Achieved

Delivery		reducing error-prone analysis and accelerating management decision cycles.
Cost	Footprint	Quantified the true cost of overstock carrying fees and lost stockout revenue,
Transparency		providing clear financial targets for process improvement initiatives.
Stakeholder Reporting		Automated the extraction and modeling pipeline, saving substantial manual data
Efficiency		assembly hours weekly and freeing analytical teams to focus on strategic insights.

The project has fundamentally transformed the organization's approach to supply chain management, establishing a data-driven culture that prioritizes proactive optimization over reactive problem-solving.

15. Technical Achievements (Resume Ready Portfolio Highlights)

This project represents a significant technical achievement, demonstrating advanced capabilities in data engineering, analytical modeling, and executive visualization.

- Engineered an enterprise-grade demand intelligence platform by developing a data pipeline spanning automated Python ETL cleaning scripts and interactive Power BI star schema relationships.
- Programmed 20+ advanced DAX calculations (such as Forecast Confidence and Estimated Lost Revenue) to translate mathematical variances into direct financial impact indicators.
- Designed a multi-tiered executive visual dashboard suite incorporating deep drill-through architecture, providing seamless monitoring from macro metrics down to daily transactional rows.
- Deployed an automated risk scoring model and prescriptive recommendation matrix, reducing manual diagnostic times and accelerating corporate resource alignment workflows.
- Led multi-regional variance analysis across multiple product categories, pinpointing localized supply constraints to recover working capital and optimize warehouse holding footprints.

16. Future Technology Roadmap & Project Conclusion

To maintain competitive positioning, the corporate analytics platform should scale towards predictive automation over three core phases:

1. **Integrating Machine Learning-based forecasting models** (such as Prophet or XGBoost) to capture advanced seasonal demand signals and complex market dynamics.
2. **Migrating batch-processed data models to cloud data lakes** (such as Azure Synapse or AWS Snowflake) linked to live ERP platforms to unlock real-time inventory monitoring and continuous intelligence.
3. **Developing dynamic what-if simulation sandboxes** using advanced DAX parameters to model supply chain adjustments under severe macro shocks or extreme market shifts, enabling robust scenario planning.

In conclusion, this project demonstrates how traditional demand forecasting can transition from an isolated mathematical exercise into an integrated enterprise decision-support system. By connecting data engineering pipelines with a clear, interactive visual narrative, the deployed architecture gives corporate stakeholders the precise insight required to reduce financial risk, optimize warehouse capital, and

Submitted by : Ankush Ghosh, PGDM 2026 – 2028, IMI Kolkata

protect net operational margins. The foundation established by this initiative positions the organization for sustained competitive advantage in an increasingly volatile global marketplace.

Appendix A – Key Python Engineering Components

To ensure scalability, reproducibility, and enterprise-grade data quality, the project leveraged Python-based data engineering workflows for ingestion, standardization, cleansing, feature engineering, and KPI generation prior to Power BI deployment.

A.1 Multi-Source Data Ingestion & Encoding Handling

```
uploaded = files.upload()
```

```
for file in file_list:
```

```
try:
```

```
df_temp = pd.read_csv(file, low_memory=False)
```

```
except UnicodeDecodeError:
```

```
try:
```

```
df_temp = pd.read_csv(file, encoding='latin1')
```

```
except:
```

```
df_temp = pd.read_csv(file, encoding='cp1252')
```

Business Value: This automation enables the framework to accommodate heterogeneous operational datasets without requiring manual intervention or file reformatting.

A.2 Schema Standardization & Data Harmonization

```
if 'store' in cols:
```

```
df['Region'] = df[lower_case_cols_map.get('store')]
```

```
elif 'state' in cols:
```

```
df['Region'] = df[lower_case_cols_map.get('state')]
```

Business Value: Creates a standardized enterprise dataset suitable for downstream analytics and reporting.

A.3 Missing Value Treatment

```
df['Units_Sold'] = df['Units_Sold'].ffill()
```

Business Value: Maintains data completeness while minimizing distortion of historical demand trends.

A.4 Outlier Detection using IQR Method

```
Q1 = df['Units_Sold'].quantile(0.25)
```

```
Q3 = df['Units_Sold'].quantile(0.75)
```

```
IQR = Q3 - Q1
```

```
df = df[
```

```
(df['Units_Sold'] > Q1 - 1.5 * IQR) &
```

```
(df['Units_Sold'] < Q3 + 1.5 * IQR)
```

```
]
```

Business Value: Reduces forecasting distortion caused by abnormal transactional spikes and data-entry anomalies.

A.5 Seasonality Modeling

```
seasonality = 1 + 0.2 * np.sin(
2 * np.pi * df['Month'] / 12
)
```

```
df['Units_Sold'] *= seasonality
```

Business Value: Introduces cyclical demand patterns commonly observed in real-world retail and supply-chain environments.

A.6 Regional Demand Adjustment

```
region_bias = {
'North': 1.1,
'South': 1.0,
'West': 1.05,
'East': 0.95
}
```

```
df['Units_Sold'] *= df['Region'].map(region_bias)
```

Business Value: Simulates geographic demand heterogeneity and supports regional planning analysis.

A.7 Forecast Generation Engine

```
bias = np.random.normal(1.0, std_dev)
bias = np.clip(bias, 0.7, 1.3)
```

```
df['Forecasted_Units'] = (
df['Units_Sold'] * bias
)
```

Business Value: Creates realistic forecast deviations for evaluating forecasting performance and operational impacts.

A.8 Inventory Allocation Logic

```
df['Inventory_Level'] = np.maximum(
df['Forecasted_Units'] * confidence,
df['Units_Sold'] *
np.random.uniform(0.95, 1.15, len(df))
)
```

Business Value: Simulates practical inventory management decisions and stock protection mechanisms.

A.9 Forecast Confidence Metric

```
df['Forecast_Confidence'] = 1 - (
abs(
df['Units_Sold']
```

Submitted by : Ankush Ghosh, PGDM 2026 – 2028, IMI Kolkata

```
- df['Forecasted_Units']  
)/ df['Units_Sold']  
)
```

Business Value: Provides an intuitive KPI for evaluating demand-planning effectiveness.

A.10 Stock Risk Identification

```
df['Stock_Risk'] = np.where(  
(df['Stockout_Flag'] == 1) |  
(df['Forecast_Confidence'] < 0.7),  
1,  
0  
)
```

Business Value: Enables proactive identification of high-risk supply-chain conditions and supports decision intelligence.